

# Homework 05: Applied Predictive Modeling

Dongzhe Li

104434089

## Exercise 6.1

**(a) Start R and use these commands to load the data:**

*library(caret)*

*data(tecator)*


**(b) In this example the predictors are the measurements at the individual frequencies. Because the frequencies lie in a systematic order (850–1,050nm), the predictors have a high degree of correlation. Hence, the data lie in a smaller dimension than the total number of predictors (215). Use PCA to determine the effective dimension of these data. What is the effective dimension?**

First, apply PCA to the data

*pc=prcomp(absorp, scale = TRUE)*

Then we can get the variance of each component and calculate the percentage of each variance. The first 6 percentages of variance are listed below:

98.626192582	0.969705229	0.279324276
0.114429868	0.006460911	0.002624591



From the figure we can see that the first component account for most of the information, so there is only one effective dimension based on the linear analysis.

**(c) Split the data into a training and a test set, pre-process the data, and build each variety of models described in this chapter. For those models with tuning parameters, what are the optimal values of the tuning parameter(s)?**

The training and test set is built below

```
Training=createDataPartition(endpoints[, 3], p = 0.75, list= FALSE)
```

```
absorbtrain=absorp[Training,]
```

```
absorbtest=absorp[-Training,]
```

```
proteintrain =endpoints[ Training, 3]
```

```
proteintest = endpoints[-Training,3]
```

In this case we only choose the third column of endpoints----the protein as the predictor. Build the linear model

```
lm=train(x = absorbtrain, y = proteintrain, method = "lm", trControl  
= trainControl(method = "repeatedcv", repeats = 5))
```

The summary is printed below:

### Linear Regression

163 samples

100 predictors

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 147, 147, 146, 147, 147, 147, ...

### Resampling results

RMSE	Rsquared	RMSE SD	Rsquared SD
1.52	0.819	0.781	0.175

The RMSE is 1.52

The Robust linear regression "rlm" model is built below, however, rlm does not allow the covariance matrix of the predictors to be singular to ensure that predictors are not singular, we will pre-process the predictors using PCA

```
rlm = train( x=absorbtrain, y=proteintrain, method="rlm",  
preProcess=c("pca"),  
trControl=trainControl(method="repeatedcv",repeats=5) )
```

### Robust Linear Model

163 samples

100 predictors

Pre-processing: principal component signal extraction, scaled, centered

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 146, 148, 147, 146, 147, 147, ...

Resampling results

RMSE	Rsquared	RMSE SD	Rsquared SD
2.6	0.275	0.349	0.17

The PLS model:

Partial Least Squares

163 samples

100 predictors

Pre-processing: centered, scaled

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 146, 148, 147, 146, 147, 147, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	RMSE SD	Rsquared SD
1	2.93	0.0975	0.274	0.115
2	2.29	0.44	0.333	0.163
3	1.75	0.664	0.404	0.167
4	1.61	0.712	0.401	0.159
5	1.18	0.854	0.177	0.0467
6	1.11	0.871	0.166	0.0437
7	1.08	0.876	0.163	0.048
8	0.95	0.906	0.167	0.0412
9	0.911	0.915	0.167	0.0389
10	0.86	0.924	0.166	0.0329
11	0.792	0.935	0.136	0.0252
12	0.733	0.947	0.154	0.0226
13	0.715	0.949	0.176	0.0236
14	0.694	0.953	0.159	0.0233
15	0.72	0.95	0.174	0.0233
16	0.804	0.934	0.277	0.0448
17	0.876	0.916	0.371	0.0818
18	0.927	0.902	0.425	0.101

19	0.95	0.894	0.472	0.117
20	0.941	0.894	0.483	0.121
21	0.921	0.899	0.459	0.112
22	1	0.877	0.566	0.161
23	1.04	0.869	0.595	0.172
24	1.04	0.872	0.564	0.152
25	1.08	0.862	0.593	0.165

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was  $ncomp = 14$ .

And the PCR :

### Principal Component Analysis

163 samples

100 predictors

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 145, 147, 148, 147, 147, 147, ...

Resampling results across tuning parameters:

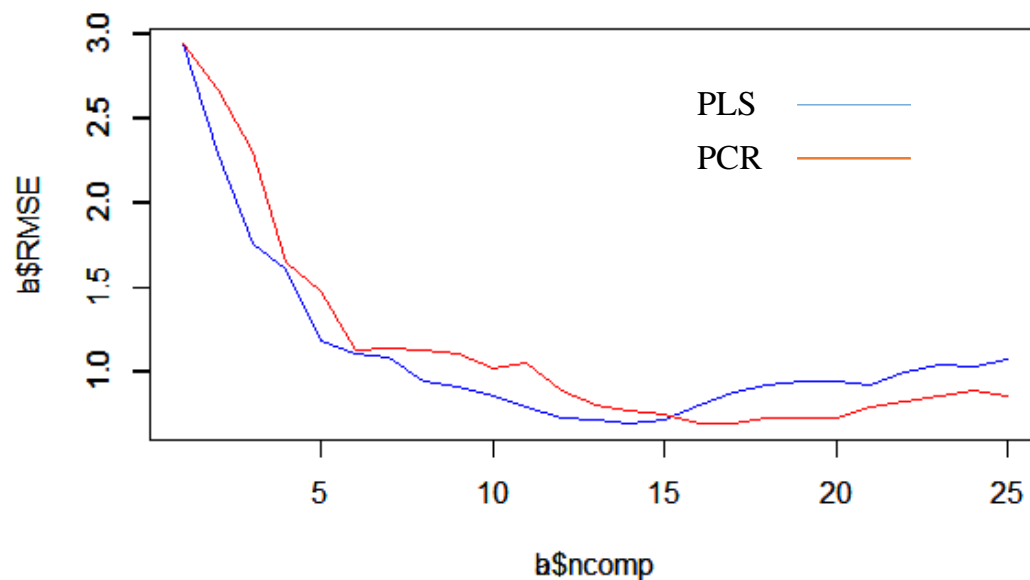
ncomp	RMSE	Rsquared	RMSE SD	Rsquared SD
1	2.94	0.106	0.274	0.113
2	2.68	0.234	0.273	0.165
3	2.3	0.441	0.369	0.18
4	1.65	0.697	0.43	0.176
5	1.47	0.766	0.286	0.102
6	1.13	0.869	0.174	0.0413
7	1.13	0.868	0.174	0.0408
8	1.12	0.868	0.168	0.04
9	1.1	0.876	0.147	0.0333
10	1.02	0.898	0.158	0.0328
11	1.05	0.891	0.174	0.0371
12	0.889	0.923	0.117	0.0231
13	0.795	0.937	0.145	0.0233
14	0.765	0.942	0.139	0.0228
15	0.744	0.944	0.137	0.0211
16	0.684	0.953	0.136	0.0214
17	0.691	0.953	0.142	0.0213

18	0.721	0.951	0.159	0.0208
19	0.718	0.951	0.18	0.0252
20	0.722	0.95	0.19	0.0265
21	0.782	0.94	0.247	0.0373
22	0.819	0.932	0.296	0.0551
23	0.853	0.925	0.324	0.0594
24	0.882	0.918	0.362	0.0683
25	0.853	0.922	0.341	0.07

RMSE was used to select the optimal model using the smallest value.


The final value used for the model was `ncomp = 16`.

These two models have similar RMSE but the PLS requires less components.



At last we build the elastic net model:

```
ENet = train(x = absorbtrain, y = proteintrain, method = "enet", trControl
= trainControl(method = "repeatedcv", repeats = 5), preProcess =
c("center", "scale"), tuneGrid = expand.grid(lambda = c(0, .001, .01, .1, 1),
fraction = seq(0.05, 1, length = 20)))
```



**(d) Which model has the best predictive ability? Is any model significantly better or worse than the others?**

From the RMSE result of different model, we find that PLS model has better performance since it is especially suited to handling highly correlated data. And linear model has the worst performance overall.

**(e) Explain which model you would use for predicting the fat content of a sample.**

PLS model can be used to predict the fat content as predicting protein in the previous problems.

## Exercise 6.2

(a) Start R and use these commands to load the data:

```
library (AppliedPredictiveModeling)
data (permeability)
```

(b) The fingerprint predictors indicate the presence or absence of substructures of a molecule and are often sparse meaning that relatively few of the molecules contain each substructure. Filter out the predictors that have low frequencies using the **nearZeroVar** function from the caret package. How many predictors are left for modeling?

```
zero = nearZeroVar( fingerprints )
fingerprints = fingerprints[, -zero]
```

There are 719 near-zero variance fingerprints, leaving 388 left for modeling  
There are  $165 \times 1107 = 182655$  elements in the previous matrix and after removing 0, there are only  $165 \times 388 = 64020$  elements.

(c) Split the data into a training and a test set, pre-process the data, and tune a partial least squares model. How many latent variables are optimal and what is the corresponding resampled estimate of **R<sup>2</sup>**?



```
Training = createDataPartition( permeability, p=0.75 )
```

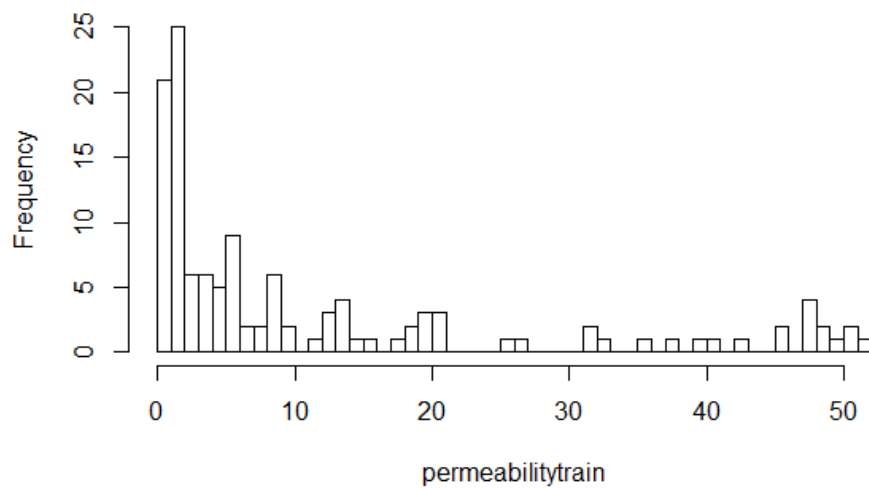
Set p=0.75 and

```
fingerprintstrain <- fingerprintsnonzero [Training$Resample1,]
```


```
permeabilitytrain <- permeability[Training$Resample1,]
```

Create the training dataset for PLS model, however, the permeability is not a symmetric distribution. We also need to pre-process the permeability before building the model

**Histogram of permeabilitytrain**



**Histogram of log2(permeabilitytrain)**



I choose log2 to modify the data and build the model.

The result of PLS model (tunelength=25) is listed below:

### Partial Least Squares

125 samples

388 predictors

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 0.75%)

Summary of sample sizes: 96, 96, 96, 96, 96, 96, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	RMSE SD	Rsquared SD
1	1.93	0.281	0.272	0.145
2	1.71	0.436	0.302	0.142
3	1.66	0.471	0.279	0.123
4	1.67	0.475	0.29	0.119
5	1.68	0.478	0.251	0.111
6	1.69	0.485	0.272	0.114
7	1.68	0.493	0.276	0.116
8	1.65	0.513	0.266	0.11
9	1.63	0.528	0.245	0.101
10	1.64	0.537	0.233	0.0949
11	1.69	0.523	0.239	0.0982
12	1.72	0.513	0.237	0.0943
13	1.75	0.503	0.246	0.104
14	1.81	0.479	0.243	0.108
15	1.84	0.471	0.252	0.112

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was ncomp = 9.

when n=10, the Rsquared reach the maximum value 0.537

**(d) Predict the response for the test set. What is the test set estimate of  $R^2$ ?**

```
ypredict = predict( plsTune, newdata=fingerprintstest )
```

And the R-squared is 0.56 which is not very far from the real value.

**(e) Try building other models discussed in this chapter. Do any have better predictive performance?**

For different models, I test lm, rlm, PCR, Enet.

LM:

```
lm = train( fingerprintstrain, permeabilitytrain, method="lm",  
preProcess=c("center","scale"),  
trControl=trainControl(method="repeatedcv",repeats=5) )
```

## Linear Regression

125 samples

388 predictors

Pre-processing: centered, scaled

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 113, 111, 112, 112, 113, 113, ...

## Resampling results

RMSE	Rsquared	RMSE SD	Rsquared SD
32.5	0.192	11.1	0.193

rlm

### Robust Linear Model

125 samples

388 predictors

Pre-processing: principal component signal extraction, scaled, centered

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 113, 111, 112, 112, 113, 113, ...

### Resampling results

RMSE	Rsquared	RMSE SD	Rsquared SD
11.7	0.517	3.57	0.235

PCR

### Principal Component Analysis

125 samples

388 predictors

Pre-processing: principal component signal extraction, scaled, centered

Resampling: Repeated Train/Test Splits Estimated (25 reps, 0.75%)

Summary of sample sizes: 96, 96, 96, 96, 96, 96, ...

### Resampling results across tuning parameters:


ncomp	RMSE	Rsquared	RMSE SD	Rsquared SD
1	16	0.0599	1.47	0.0787
2	16	0.0597	1.47	0.0785
3	14.8	0.196	0.0232	0.229
4	12.5	0.401	0.445	0.164
5	11.9	0.466	0.652	0.127
6	11.9	0.461	0.657	0.127
7	12.3	0.416	0.538	0.125

8	12.3	0.427	0.791	0.132
9	12.7	0.393	0.246	0.177
10	12.4	0.432	0.697	0.125
11	12	0.47	0.293	0.153
12	11.9	0.471	0.238	0.152
13	12	0.463	0.257	0.152
14	12.1	0.458	0.0867	0.164
15	12	0.461	0.0297	0.168
16	12.3	0.44	0.14	0.163
17	12.2	0.462	0.0875	0.129
18	11.9	0.461	0.144	0.155
19	11.9	0.457	0.289	0.14
20	11.9	0.461	0.127	0.146
21	11.6	0.481	0.374	0.18
22	11.5	0.486	0.315	0.175
23	11.5	0.488	0.328	0.178
24	11.6	0.493	0.326	0.143
25	12	0.48	0.18	0.157

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was  $n_{\text{comp}} = 23$

Enet



**(f) Would you recommend any of your models to replace the permeability laboratory experiment?**

From the result above, the elastic net solution the best penalized model choice for this problem. It has the lowest RMSE value.

## Appendix

Q1

```
library(caret)
data(tecator)
```

```
#calculate the mean and variance of each column of endpoints
apply(absorp, 2, mean)
apply(absorp, 2, var)
```

```
pc=prcomp(absorp, scale = TRUE)
```

```
print( pc$center )
print( pc$scale )
```

```
biplot(pc, scale = 0)
```

```
#calculate the largest 6 eigenvalues percentage
pc_var = pc$sdev^2
pve = pc_var/sum(pc_var)*100
head(pve,10)
plot( pve, xlim=c(0,40), type='b', pch=16, xlab='number of components',
      ylab='percent of total variation' )
# 98.626192582  0.969705229  0.279324276  0.114429868  0.006460911
0.002624591
#only the first dimension is effective
#we use the third attribute in the endpoints
set.seed(0)
```

```
Training=createDataPartition(endpoints[, 3], p = 0.75, list= FALSE)
```

```
absorbtrain=absorp[Training,]
```

```

absorbtest=absorp[-Training,]

proteintrain =endpoints[ Training, 3]
proteintest = endpoints[-Training,3]

control = trainControl(method = "repeatedcv", repeats = 5)

#linear model
set.seed(1)
lm=train(x = absorbtrain, y = proteintrain, method = "lm", trControl =
trainControl(method = "repeatedcv", repeats = 5))
mean(proteintrain)

#rlm
rlm = train( x = absorbtrain, y = proteintrain, method="rlm",
preProcess=c("pca"), trControl=trainControl(method="repeatedcv",repeats=5) )

#pcr
set.seed(1)
PCR=train(x = absorbtrain, y = proteintrain, method = "pcr",
preProcess=c("pca"),trControl = control, tuneLength = 25)
#pls
set.seed(1)
PLS=train(x = absorbtrain, y = proteintrain, method =
"pls",trControl=trainControl(method="repeatedcv",repeats=5), preProcess =
c("center", "scale"),tuneLength = 25)
#PCR
PCR <- train(x = absorbtrain, y = proteintrain,method = "pcr", trControl =
trainControl(method="repeatedcv",repeats=5), tuneLength = 25)

comps <- rbind(PLS$results, PCR$results)
comps$Model <- rep(c("PLS", "PCR"), each = 25)
a=comps[1:25,1:6]
b=comps[26:50,1:6]

plot(a$ncomp,a$RMSE,col="blue",type="l")
par(new=TRUE)
plot(b$ncomp,b$RMSE,col="red",type="l")
legend(col=c("red","blue"))
set.seed(1)
#enet

```

```
ENet = train(x = absorbtrain, y = proteintrain, method = "enet", trControl =
trainControl(method="repeatedcv", repeats=5), preProcess = c("center",
"scale"), tuneGrid = expand.grid(lambda = c(0, .001, .01, .1, 1), fraction =
seq(0.05, 1, length = 20)))
plot(ENet)
```

Q2

```
library(caret)
library(AppliedPredictiveModeling)
data(permeability)
```

```
zero = nearZeroVar( fingerprints )
fingerprintsnonzero = fingerprints[,-zero]
```

```
Training = createDataPartition( permeability, p=0.75 )
fingerprintstrain = fingerprintsnonzero [Training$Resample1,]
permeabilitytrain =permeability[Training$Resample1,]
```

```
fingerprintstest=fingerprintsnonzero [-Training$Resample1,]
permeabilitytest=permeability[-Training$Resample1,]
```

```
hist(permeabilitytrain,breaks=50)
hist(log2(permeabilitytrain),breaks=50)
set.seed(3)
control=trainControl(method = "LGOCV")
plsTune =train(x = fingerprintstrain, y =log10(permeabilitytrain) , method =
"pls",tuneGrid = expand.grid(ncomp = 1:15),trControl = control)
plsTune
```

```
ypredict = predict( plsTune, newdata=fingerprintstest )
rsquared_pls = cor(ypredict,permeabilitytest,method="pearson")^2
```

```
set.seed(3)
lm = train( fingerprintstrain, permeabilitytrain, method="lm",
preProcess=c("center","scale"),
trControl=trainControl(method="repeatedcv",repeats=5) )
ypredict_lm = predict( lm, newdata=fingerprintstest )
r2_lm = cor(ypredict_lm,permeabilitytest,method="pearson")^2
```



```
set.seed(3)
rlm = train( fingerprintstrain, permeabilitytrain, method="rlm",
preProcess=c("pca"), trControl=trainControl(method="repeatedcv",repeats=5) )
ypredict_rlm = predict( rlm, newdata=fingerprintstest )
r2_rlm = cor(ypredict_rlm,permeabilitytest,method="pearson")^2
```

```
set.seed(3)
PCR=train(fingerprintstrain, permeabilitytrain, method = "pcr",
preProcess=c("pca"),trControl = control, tuneLength = 25)
```

```
set.seed(3)
```

```
enet = train( fingerprintstrain, permeabilitytrain, method="enet",
tuneGrid = expand.grid(lambda = c(0, .05, .1), fraction = seq(0.05, 1, length =
20)), trControl=trainControl(method="repeatedcv",repeats=5) )
plot(enet)
```